

# Automated detection of Hainan gibbon calls for passive acoustic monitoring

Emmanuel Dufourq<sup>\*1,2</sup>, Ian Durbach<sup>3,4,1</sup>, James P. Hansford<sup>5, 6</sup>, Amanda Hoepfner<sup>8</sup>, Heidi Ma<sup>5</sup>, Jessica V. Bryant<sup>10</sup>, Christina S. Stender<sup>7</sup>, Wenying Li<sup>9</sup>, Zhiwei Liu<sup>9</sup>, Qing Chen<sup>9</sup>, Zhaoli Zhou<sup>9</sup>, and Samuel T. Turvey<sup>5</sup>

<sup>1</sup>African Institute for Mathematical Sciences, South Africa

<sup>2</sup>Stellenbosch University, South Africa

<sup>3</sup>Centre for Research into Ecological and Environmental Modelling, University of St Andrews, UK

<sup>4</sup>Centre for Statistics in Ecology, the Environment, and Conservation, University of Cape Town, South Africa

<sup>5</sup>Institute of Zoology, Zoological Society of London, Regent's Park, London NW1 4RY, UK

<sup>6</sup>Department of Biological Sciences, Northern Illinois University, DeKalb, Illinois 60115, USA

<sup>7</sup>Living Collections, Zoological Society of London, Regent's Park, London NW1 4RY, UK

<sup>8</sup>School of Biological Sciences, University of Utah, Salt Lake City, UT 84112, USA

<sup>9</sup>Bawangling National Nature Reserve, Changjiang Li Autonomous County, Hainan, China

<sup>10</sup>Department of Life Sciences, University of Roehampton, London SW15 4JD, UK

September 14, 2020

## 1 Abstract

1. Extracting species calls from passive acoustic recordings is a common preliminary step to ecological analysis. For many species, particularly those occupying noisy, acoustically variable habitats, the call extraction process continues to be largely manual, a time-consuming and increasingly unsustainable process. Deep neural networks have been shown to offer excellent performance across a range of acoustic classification applications, but are relatively underused in ecology.
2. We describe the steps involved in developing an automated classifier for a passive acoustic monitoring project, using the identification of calls of the Hainan gibbon (*Nomascus hainanus*), one of the world's rarest mammal species, as a case study. This includes pre-processing - selecting a temporal resolution, windowing and annotation; data augmentation; processing - choosing and fitting appropriate neural network models; and postprocessing - linking model predictions to replace, or more likely facilitate, manual labelling.

---

\*Corresponding author: [edufourq@gmail.com](mailto:edufourq@gmail.com)

30 3. Our best model converted acoustic recordings into spectrogram images on the mel frequency  
31 scale, using these to train a convolutional neural network. Model predictions were highly  
32 accurate, with per-second false positive and false negative rates of 1.5% and 22.3%. Nearly  
33 all false negatives were at the fringes of calls, adjacent to segments where the call was correctly  
34 identified, so that very few calls were missed altogether. A postprocessing step identifying  
35 intervals of repeated calling reduced an eight-hour recording to, on average, 22 minutes for  
36 manual processing, and did not miss any calling bouts over 72 hours of test recordings.  
37 Gibbon calling bouts were detected regularly in multi-month recordings from all selected  
38 survey points within Bawangling National Nature Reserve, Hainan.

39 4. We demonstrate that passive acoustic monitoring incorporating an automated classifier rep-  
40 represents an effective tool for remote detection of one of the world’s rarest and most threatened  
41 species. Our study highlights the viability of using neural networks to automate or greatly  
42 assist the manual labelling of data collected by passive acoustic monitoring projects. We em-  
43 phasise that model development and implementation be informed and guided by ecological  
44 objectives, and increase accessibility of these tools with a series of notebooks that allow users  
45 to build and deploy their own acoustic classifiers.

46 **Keywords:** bioacoustics, passive acoustic monitoring, species identification, deep learning, con-  
47 volutional neural networks, Hainan gibbons.

## 48 2 Introduction

49 Deep learning holds enormous promise for automating the labelling of bioacoustic data. The num-  
50 ber of applications is growing (Christin, Hervet, & Lecomte, 2019), but the majority of datasets are  
51 still labelled manually (Fairbrass et al., 2019; Kiskin et al., 2020; Pamula, Pocha, & Klaczynski,  
52 2019), even as the rate of data collection makes this approach increasingly unsustainable. The  
53 mismatch between the potential of deep learning approaches and their actual uptake among prac-  
54 titioners occurs because getting models to perform as well as an experienced human is difficult.  
55 Human-like performance usually requires substantial amounts of training data or relatively sta-  
56 ble background environments, conditions that are often absent in ecological applications. Model  
57 tuning and data manipulation is often required, and while guidelines are emerging (Patterson &  
58 Gibson, 2017; Stowell, Wood, Pamula, Stylianou, & Glotin, 2019), these can, with some justifica-  
59 tion, appear subjective and case specific. A lack of computing resources and user-friendly software  
60 can also be a barrier to entry. Case studies reporting successful applications play an important  
61 role in developing and disseminating best practices, and in discriminating between those tasks that  
62 current deep learning methods are able to automate and those they cannot. Previous applications  
63 have used convolutional neural networks (CNNs; LeCun, Bengio, and Hinton (2015)) to identify

64 various bird (Grill & Schlüter, 2017; Kahl et al., 2017; Stowell, Wood, et al., 2019) and whale  
65 species (Bergler et al., 2019; Bermant, Bronstein, Wood, Gero, & Gruber, 2019; Jiang et al., 2019;  
66 Shiu et al., 2020), bees (Kulyukin, Mukherjee, & Amlathe, 2018; Nolasco et al., 2019), as well as  
67 anomalous acoustic events in soundscapes (Sethi et al., 2020). These have shown, for example,  
68 that a generally good approach is to represent data as spectrograms and treat the problem as  
69 an image classification one, as well as providing specialised approaches for data augmentation on  
70 spectrogram inputs, such as pitch and time shifting and introducing background noise (Bergler et  
71 al., 2019; Sprengel, Jaggi, Kilcher, & Hofmann, 2016).

72 Despite this, no studies report the process of applying deep learning within the scope of a  
73 typical acoustic monitoring project designed to answer a well-defined research question. Most  
74 applications are either smaller – using data collected for the purpose of testing a deep learning  
75 approach, and often written for a machine learning rather than ecological audience (e.g. Kiskin  
76 et al., 2020; Kulyukin et al., 2018); or larger – aggregating datasets across several independent  
77 studies to investigate if models generalise (Bergler et al., 2019; Shiu et al., 2020; Stowell, Wood,  
78 et al., 2019) – than most monitoring projects. In this paper we address this gap, describing the  
79 development of a classifier for identifying Hainan gibbon (*Nomascus hainanus*) calls in passive  
80 acoustic recordings collected as part of a long-term monitoring project, with the aim of providing  
81 practitioners with a realistic and relatable idea of the process, and modelling choices, involved, as  
82 well as guidelines for these choices.

83 The Hainan gibbon is the world's rarest primate and one of the world's rarest mammals, with  
84 only a single population of about 30 individuals surviving in Bawangling National Nature Reserve  
85 (BNNR), Hainan, China (Chan, Fellowes, Geissmann, & Zhang, 2005; Liu, Ma, Cheyne, & Turvey,  
86 2020; S. Turvey et al., 2015). Improved monitoring of this population using novel methods, to un-  
87 derstand factors affecting successful dispersal, breeding group formation and colonization of new  
88 habitat, has been identified as an urgent short-term conservation goal for the species (S. Turvey  
89 et al., 2015; Zhang et al., 2020). Gibbons call regularly to advertise territory and maintain group  
90 cohesiveness against rivals, using a complex structure consisting of short individual vocal sylla-  
91 bles or "notes" of ca. 0.2–2.75s assembled together into longer "phrases" consisting of one to six  
92 notes, which are themselves organised into "songs" of several minutes (Deng, Zhou, & Yang, 2014).  
93 Gibbon population surveys are usually conducted by detecting this daily song using a fixed-point  
94 count survey method, whereby researchers listen opportunistically for calls at elevated listening  
95 posts (Brockelman & Srikosamatara, 1993; Kidney et al., 2016). However, this traditional moni-  
96 toring approach is labour-intensive and is only conducted for discrete survey periods. Gibbons are  
97 therefore prime candidates for passive acoustic monitoring and recent studies have used data col-  
98 lected in this way to model occupancy (Vu & Tran, 2019) and to discriminate between individuals  
99 using spectral features (Clink, Crofoot, & Marshall, 2019; Zhou et al., 2019). All of these studies,  
100 however, have relied on an initial manual extraction of calls.

101 In order to develop a continuous monitoring protocol for Hainan gibbons we conducted long-  
 102 term passive acoustic monitoring and developed an automated classifier able to identify whether  
 103 gibbons were calling in the vicinity of a particular recorder, with the aim of establishing whether  
 104 the area proximal to the recorder was occupied that day. It was therefore important to be able  
 105 to detect individual gibbon calling bouts, but not necessarily to be able to discriminate every  
 106 phrase made during the bout. We address issues that are important to the overall usefulness of a  
 107 classifier, including deciding how much data to manually label, data augmentation, operationally  
 108 meaningful definitions of classifier success, and the development of user-friendly software. Our  
 109 study provides an effective new monitoring method for the world's rarest primate, and also has  
 110 wider applicability for applying deep learning to develop passive acoustic monitoring frameworks  
 111 for other conservation-priority loud-call species such as cetaceans, elephants, or other primates  
 112 (Crunchant, Borchers, Kuhl, & Piel, 2020).

## 113 3 Materials and Methods

### 114 3.1 Data collection

115 Eight Song Meter SM3 recorders (Wildlife Acoustics, Maynard, Massachusetts) were used to collect  
 116 acoustic data from 1 March to 20 August 2016 within BNNR. Recorders were attached to trees at  
 117 a height of approximately 1.5m in tropical evergreen forest. Four recorders were situated within  
 118 the known home ranges of the four Hainan gibbon social groups existing during the study period  
 119 (Groups A-D; see Bryant, Zeng, Hong, Chatterjee, and Turvey (2017)), three were situated at  
 120 locations intermediate between known home ranges, and a further recorder was placed in an area  
 121 where a solitary male gibbon was thought to occur (Bryant et al., 2016). They were placed at  
 122 locations that were used as regular listening posts for monitoring gibbons by reserve staff (Figure  
 123 1). The peak Hainan gibbon calling period is 06:00–07:00, with calling continuing at decreasing  
 124 regularity for several hours (Chan et al., 2005). Recorders were therefore set to record for eight  
 125 hours each day from the time of sunrise, which varied between approximately 05:00 and 06:00  
 126 during the study period. Memory cards and batteries were changed every 40 days. Devices did  
 127 not record continuously throughout the entire survey period due to logistical and technical issues;  
 128 in total, survey days per recorder varied between 79 and 129 days, and roughly 6,000 hours of  
 129 recordings were collected. The majority of recordings were made with a sampling rate of 9,600Hz  
 130 and bit depth of 16, with isolated recordings at 28,800Hz.

### 131 3.2 Data analysis

132 We manually labelled 32 eight-hour recordings by inspecting spectrograms and listening to audio  
 133 using Sonic Visualiser (Cannam, Landone, & Sandler, 2010), and recording the start and end

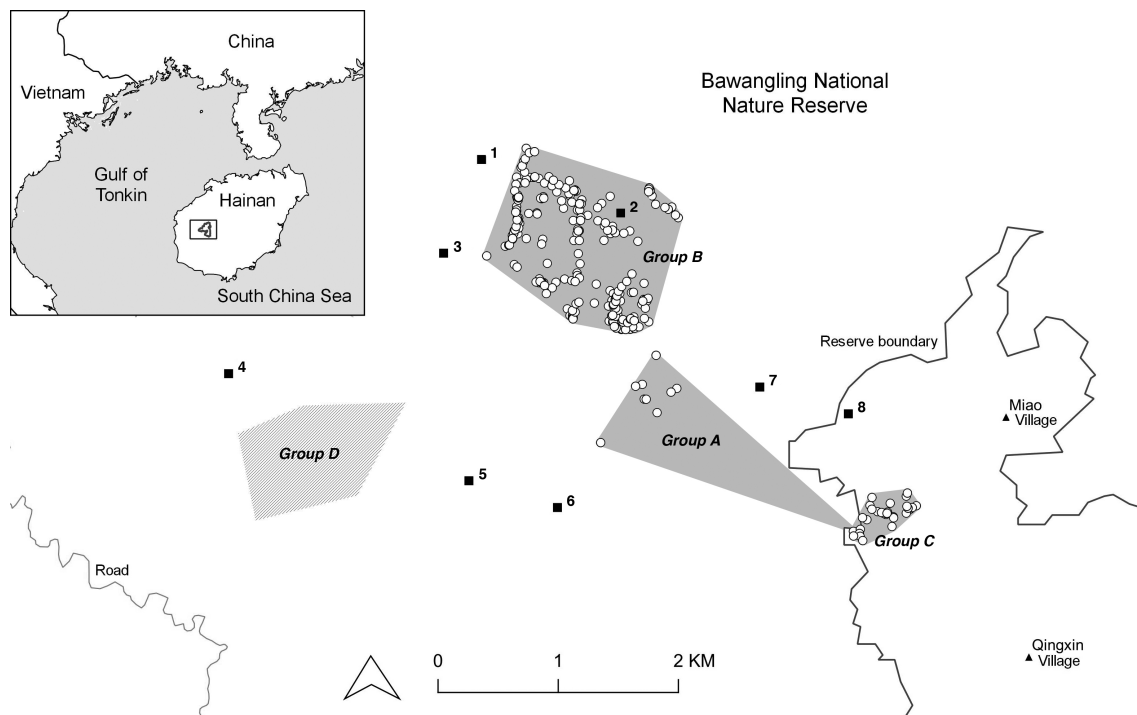


Figure 1: Locations of eight Song Meter SM3 recorders (labelled 1-8) used to detect gibbons in 2016 within Bawangling National Nature Reserve, Hainan, China, in relation to approximate distributions of four Hainan gibbon social groups (A-D). Mapped distributions of groups A-C are based on field data collected in 2010-2011 (see Bryant et al. (2017)); the groups all changed their location slightly between 2011 and 2016, but data on exact group locations in 2016 are unavailable. Approximate location of Group D indicated with hatching based on Bryant et al. (2016).

134 times, and the number of notes, of each observed gibbon phrase. This process yielded 1,246 gibbon  
135 phrases.

136 To construct the fixed-length inputs required by CNNs, we divided each eight-hour recording  
137 into segments with window length 10s and hop length 1s (starting times of consecutive 10s segments  
138 differ by 1s, Figure 2). This window length was chosen so that even the longest phrase (8s,  
139 Supplementary Material A) fits within a single segment; using a slightly longer segment length  
140 allows for potentially longer unseen phrases, and results in more positive segments after windowing.  
141 All audio was converted into mono, as done in various applications (e.g. Bergler et al., 2019; Qazi,  
142 Tabassam Nawaz, Rashid, & Habib, 2018; Stowell, Petrusková, Šálek, & Linhart, 2019). By cross-  
143 referencing the time intervals of each segment with the logged start and end times of known gibbon  
144 phrases, each segment was labelled as (a) a “presence”, if its time interval completely contained  
145 the interval of at least one labelled phrase, (b) an “absence”, if its time interval contained no part  
146 of any phrase, or (c) a “partial presence”, if its time interval intersected but did not completely  
147 contain the interval of at least one labelled phrase (Figure 2). Partial presences were excluded  
148 from further analysis.

149 Preprocessed amplitudes in each 10s segment were downsampled to 4800Hz, and the down-  
150 sampled inputs – each segment a time series of 48000 observations – used as inputs to the 1-D  
151 CNNs described in the next section. In addition, we converted each audio segment into a mel-scale

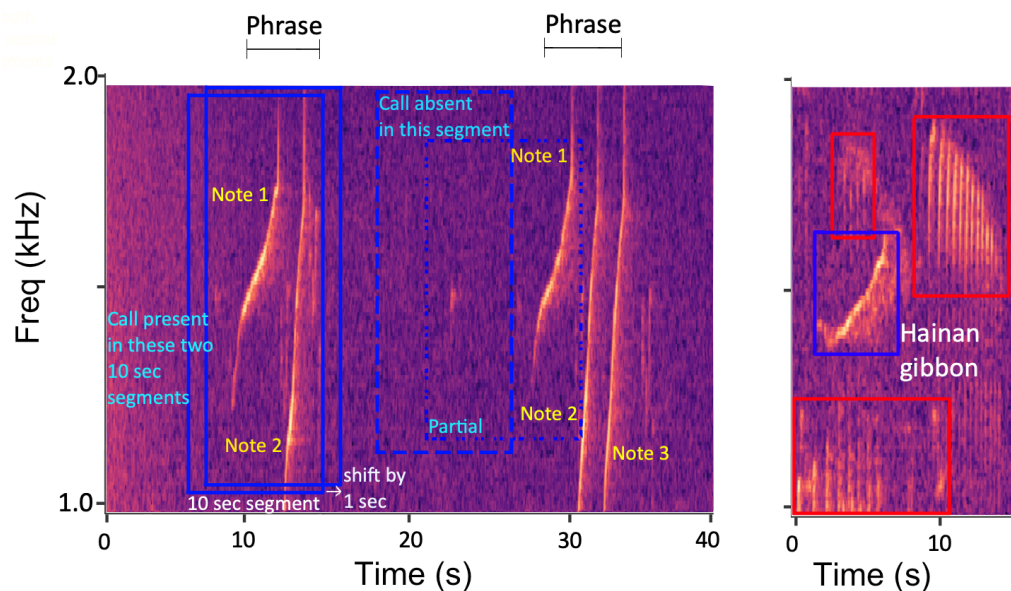


Figure 2: Hainan gibbon calls consist of a sequence of “phrases”, each phrase consisting of variable (typically, 1-6) “notes” and often with relatively large intervals between phrases. **Left:** a two-note phrase followed by a three-note phrase. A single calling bout may last anywhere from a few to dozens of minutes. Our model divides the recording interval into sliding 10s windows or “segments” (blue boxes), with 80% overlap between adjacent segments. Segments are classified as contained at least one full gibbon phrase (Present; solid line), a partial phrase (Partial; dotted line), or no part of a phrase (Absent; dashed line). Partial presences were excluded from further analysis, creating a two-class audio classification problem. **Right:** a gibbon phrase partially obscured by noisy background conditions, in this case other species calling (red boxes).

152 spectrogram (Bergler et al., 2019; Huang, Acero, & Hon, 2001), to be used as an input image to  
153 a 2-D CNN, using a window size of 1,024/9,600s, a hop size of 256/9,600s, and 128 mel frequency  
154 bins with centres uniformly spaced between 1 and 2kHz, a conservative interval following Deng et  
155 al. (2014) and our own exploratory analyses. These values for chosen on the basis of preliminary  
156 investigations, although results are not particularly sensitive to these choices. The spectrogram  
157 images had a size of  $128 \times 188$  pixels; larger image sizes can capture greater detail but typically  
158 require more network parameters and computation time to do so.

159 After processing, our dataset consisted of 5,285 segments containing at least one complete  
160 phrase. While the vast majority of segments do not contain any gibbon calls, we restricted the  
161 number of absence segments to the same number as presences, to avoid a large class imbalance.  
162 Absence segments were initially collected by randomly sampling, but we found that better results  
163 were obtained by specifically including absence segments that contained typical ambient noise, such  
164 as bird calls, rain events, and other background noises that could potentially confuse the classifier  
165 (Stowell, Petrusková, et al., 2019). Extracting these required additional manual processing of the  
166 audio data.



### 3.3 Data augmentation

Data augmentation – boosting sample sizes by adding new samples artificially created by manipulating existing ones, for example using geometric operations like translations and rotation – is commonly used to improve classifier performance, particularly when the training dataset is relatively small (Hestness et al., 2017; Sun, Shrivastava, Singh, & Gupta, 2017). We used data augmentation to create up to ten new copies of each 10s segment in both presence and absence classes. For each presence segment  $\mathbf{x}^{(pre)}$ , we randomly selected ten absence segments,  $\mathbf{x}_i^{(abs)}, i = 1, \dots, 10$ . We randomly shifted the starting time of each absence segment forward by  $0 < t_i < 9$  seconds, with the absence segment wrapping back on itself so that it remained 10s long (Figure 3c), to obtain the shifted segment  $\mathbf{x}_i^{(shift)}$ . Presence segments were not shifted, as this already occurred during the windowing process used to create the original segments. Segments contain amplitude values and thus allow for arithmetic operations to be performed on them. We blended the presence segment with each shifted segment to create augmented presence segments  $\mathbf{x}_i^{(aug)} = \alpha \mathbf{x}^{(pre)} + (1 - \alpha) \mathbf{x}_i^{(shift)}$ , where  $\alpha$  is a mixing parameter, here chosen to be 0.9 (Figure 3d). We created augmented absence segments using the same approach, i.e. combining pairs of absence segments to create a mixture of background scenes.

After augmenting the original segments, we obtained 18,992 segments (9,496 presence, 9,496 absence) from 19 recordings to train the neural networks. We randomly selected 60% of the data for training (5,697 presence, 5,697 absence) and used the remaining 40% for validation (3,799 presence, 3,799 absence). Non-augmented segments from nine separate recordings (2,231 presence, 23,689 absence) were kept aside for testing.

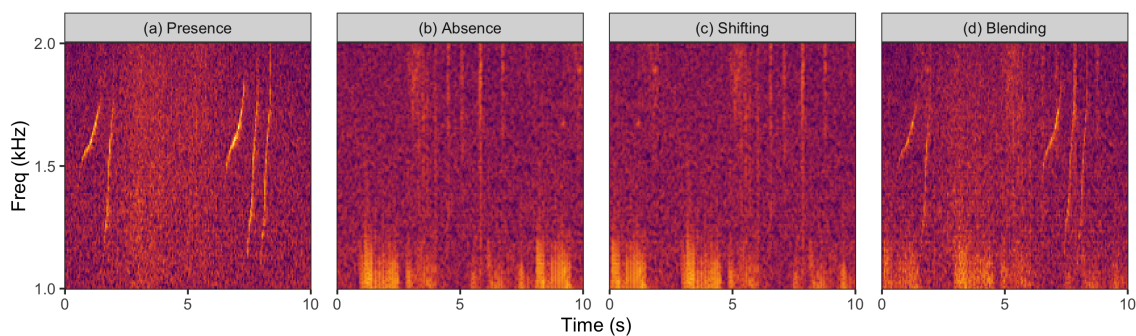


Figure 3: Data augmentation steps involve (a) selecting a presence segment containing a Hainan gibbon phrase, (b) randomly selecting a segment containing only background noise, (c) shifting the starting time of the absence segment forward by a random amount, here two seconds, and (d) blending together the presence and shifted absence segments.

### 3.4 Neural networks

We considered two kinds of CNN architectures: a 1-D CNN using preprocessed amplitudes of 10s segments as inputs, and a 2-D CNN that had inputs consisting of spectrogram images constructed from the preprocessed amplitudes. As we had relatively little training data by deep learning stan-

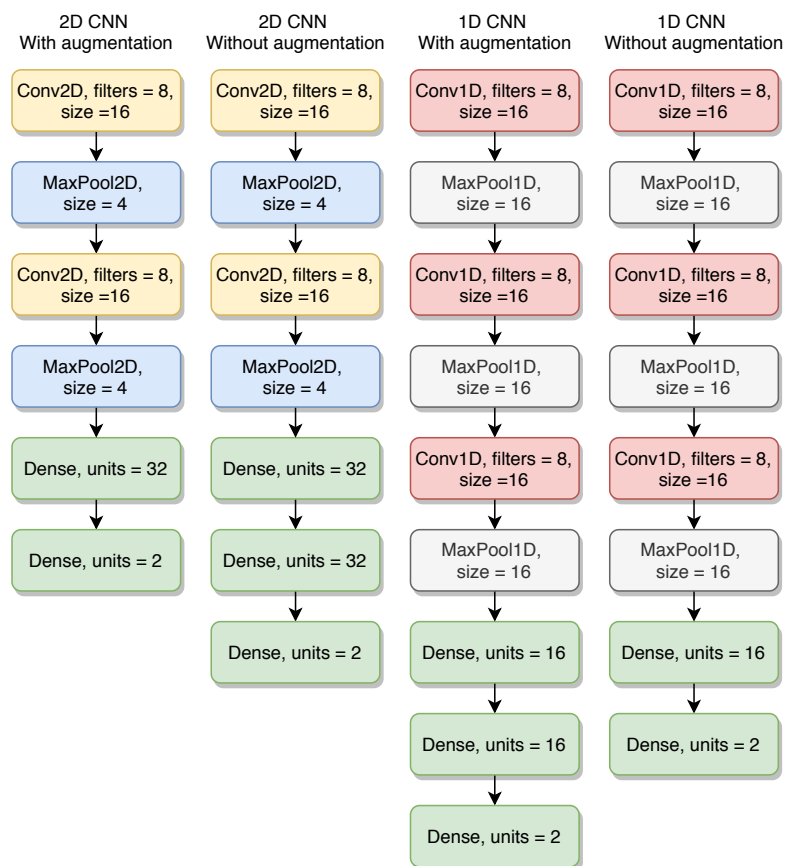


Figure 4: Best architectures for 1-D and 2-D CNNs, for both augmented and non-augmented training datasets. Selected architectures were those with intermediate numbers of free parameters, particularly for 2-D CNNs.

192 dards, we chose these networks as they use simple architectures requiring relatively few parameters.  
193 Both 1-D and 2-D CNNs use up to three convolutional layers, each followed by a max pooling layer  
194 that reduces the size of the intermediate input passed to the next layer of the network. We used  
195  $16 \times 1$  and  $16 \times 16$  convolutional kernels for 1-D and 2-D CNNs, respectively. The stack of convo-  
196 lutional layers was followed by one or two dense layers (Figure 4). The resulting model outputs a  
197 predicted probability that the input segment (1-D or 2-D) contains at least one complete gibbon  
198 phrase.

199 We chose model hyperparameters using a grid search over the number of convolutional (1, 2, 3)  
200 and dense (1, 2, 3) layers, nodes in each of the dense layers (8, 16, 32), filters in each convolutional  
201 layer (8, 16, 32), kernel size in each convolutional and max pooling layer (4, 8, 16), and dropout rate  
202 (0, 0.2, 0.4, 0.6). Each model was trained for 50 epochs using the Adam optimizer (Kingma & Ba,  
203 2014) a batch size of 8 segments, and a learning rate of 0.001. Models were evaluated based on  
204 test set accuracy (proportion of all predictions that were correct), sensitivity (proportion of true  
205 positives divided by positive examples), and specificity (proportion of true negatives divided by  
206 negative examples). Optimal thresholds for converting predicted probabilities into binary classifi-  
207 cations were those that minimized the ratio of sensitivity and false discovery rate in the validation  
208 dataset.



Models were implemented in Python 3 using the TensorFlow (Abadi et al., 2015) library with Keras (Chollet et al., 2015) for the neural network component, and the Librosa library for audio processing and spectrogram construction (McFee et al., 2020). Model training and testing was done on a machine running Ubuntu 16.04 LTS with an Intel i7-6700K CPU, 16GB of RAM, and an Nvidia GTX 1070 8GB Graphics Processing Unit. Code and analysis scripts are available online at <https://github.com/emmanueldufourq/GibbonClassifier>.

### 3.5 Post-processing

For an audio recording of arbitrary duration, our approach was to break that recording into overlapping 10s segments, and to use a trained CNN to output, for each segment starting at second  $s = 0, 1, 2, \dots$ , a predicted probability indicating the likelihood that at least one complete gibbon phrase is contained in the next ten seconds. These probabilities are based only on the acoustic content of their associated segments, and can give rise to biologically unrealistic call patterns. We used a post-processing step to remove isolated predicted presence segments which are highly likely to be false positives rather than actual calls, and to obtain start and end times for each predicted calling bout, to facilitate manual verification and support the main research objective of detecting and monitoring gibbon activity.

To do this, we formed connected components of presence segments that occur close together in time and in sufficient numbers that, given known gibbon call characteristics (i.e. song duration, inter-phrase duration), they are likely to be part of a single calling bout (Supplementary Material A). With presence segments arranged in temporal order, presence segment  $i$  is included in the same component as segment  $i-1$  if they are separated by less than 200s; otherwise segment  $i$  begins a new component. This process allocates each presence segment to exactly one component. Components were then reviewed, and any components consisting of fewer than 20 segments (equivalent to roughly four phrases of length 5s) were removed, as were any components where the average time between consecutive presence segments in the component was greater than 10s (suggesting a "chain" of isolated presence predictions, since calls usually persist over multiple consecutive segments).

The first and last presence segment in each remaining component give the start and end times of each predicted gibbon calling bout. To evaluate the post-processing step, we mimic its intended application by assuming that all predicted bouts are passed to an observer for manual processing, and that all presence segments within the bout are subsequently identified. This approach means that post-processing accuracy measures are conditional on the use of additional, error-free manual verification.

## 4 Results

Hainan gibbon calls could be detected with a high degree of accuracy. Without post-processing, nearly 80% of segments containing gibbon calls were correctly identified, with very few false positives (Table 1). Even with false negative rates of 20% very few gibbon phrases were missed altogether, because phrases occur across multiple overlapping segments and nearly all segments incorrectly identified as absences occurred at the beginning and end of a phrase, abutted by several segments where the phrase was correctly detected (Figure 5). After post-processing, fewer than 2% of all presence segments occurred outside of predicted call bouts (Table 1), and all 20 call bouts across nine test set recordings were detected, with two predicted call bouts being false positives (Supplementary Material B). In the training set, 34 of 35 call bouts were correctly recognised with 2 false positive call bouts.

CNN + Augmentation + Postprocessing	2-D Yes Yes	2-D Yes No	2-D No No	1-D Yes Yes	1-D Yes No	1-D No No
Accuracy (Test)	99.37%	97.60%	92.32%	94.30%	94.76%	94.76%
Sensitivity (Test)	98.30%	77.68%	79.65%	54.21%	40.98%	25.56%
Specificity (Test)	99.42%	98.51%	92.92%	95.96%	96.91%	97.60%
Accuracy (Train)	98.68%	97.20%	93.65%	95.14%	94.16%	93.44%
Sensitivity (Train)	94.84%	80.64%	77.85%	69.62%	53.42%	24.53%
Specificity (Train)	99.12%	98.59%	94.94%	97.66%	97.92%	99.24%
Model Parameters	23,922	23,922	24,978	2,650	2,650	2,378
Train Duration (sec)	644	643	265	628	627	117

Table 1: Average classification accuracy and parameter settings for the best 2-D and 1-D CNN models across 72 hours of test recordings (2,231 segments containing gibbon phrases, 23,689 without). Gibbon calls can be identified with very high accuracy, and performance is improved by data augmentation and a postprocessing heuristic.

The best performing approach was a 2-D CNN with both data augmentation and post-processing. Data augmentation improved specificity by 5.6%, a relative reduction in false positives of 79% but without associated relative reduction in sensitivity; post-processing further improved both sensitivity (20.6%) and specificity (0.9%, Table 1). Accuracy was substantially higher when treating the task as an image (spectrogram) classification problem than if the preprocessed acoustic data were directly used as input to a 1-D CNN. An 8 hour test file took on average 6 minutes to process of which 3 minutes 10 seconds were used for reading in the audio file and 2 minutes 42 seconds to convert to spectrograms; the remaining time was used to compute the CNN predictions.

Across the entire monitoring project, gibbon calls were detected on 71% of recording days across all locations. Gibbons were detected regularly at all locations, with recorders situated within known group or solitary home ranges detecting calls on 33–86% of recording days, and those situated between home ranges detecting calls on 46–89% of recording days. Mean durations of calling bouts per recorder varied between 24.2 and 40.8 minutes (overall mean = 29.7 minutes), with mean starting times of 06:16–07:56 am and mean finishing times of 09:12–10:15 am (Figure

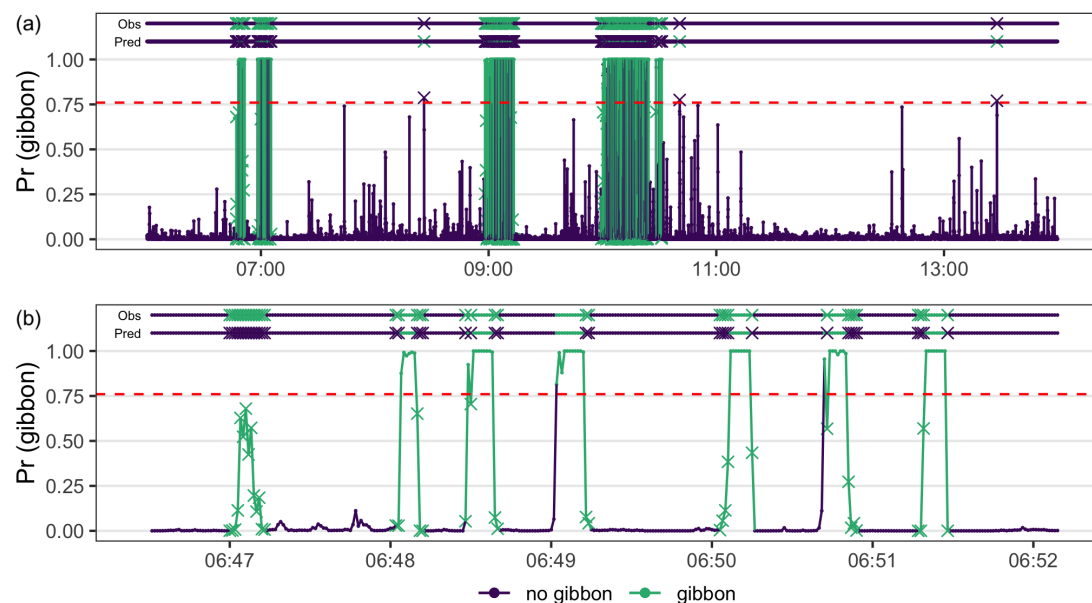


Figure 5: Per-second predicted probabilities that a gibbon phrase is contained within the next 10s of audio, over (a) an eight-hour file, (b) a five-minute window. Segments with predicted probabilities above an optimized threshold of 0.76 (red line) are classified as containing a gibbon phrase, with misclassifications denoted by crosses. Observed and predicted classes are plotted above the probabilities, using the same notation. Colour is used to denote the observed class. Most incorrect false negative classifications are at the beginning and end of phrases, separated by segments that correctly identify the call. In this way, nearly all phrases are clearly identified, and a practitioner can be pointed to those regions that contain calls.

6; Table 2). Calls were detected less frequently during the wet season (March-April) than the dry season (May-August), with inter-season differences varying substantially between locations (Supplementary Table C).

## 5 Discussion

Long-term monitoring will generate thousands of hours of recordings across multiple survey sites, and manually labelling these recordings is typically infeasible given logistical constraints. Our results demonstrate that passive acoustic monitoring incorporating an automated classifier can be an effective tool for remote detection of calling species, potentially enabling systematic monitoring whilst saving time, funds and manpower. Our approach, applied to Hainan gibbons, is general and easily extended to other calling species.

Our models allow new recordings to be classified on a per-second basis, to a high degree of accuracy. Although perhaps false negative rates of 1.7% may not be sufficiently low for full automation of Hainan gibbon call monitoring, they greatly facilitate the process of manually annotating these datasets by ruling out large portions of recordings that have a near-zero probability of containing gibbon song. In our test datasets, this reduced the amount of audio to be manually processed by 95%. Our model clearly detected all calling bouts in the test data, at the cost of two false positives.




Figure 6: Daily patterns in gibbon calling activity. The red line denotes, per 10 minutes, the proportion of recordings across all locations in which a call was detected (e.g. 05:00-05:10, 05:10-05:20, ...). The black line smooths the observed proportions using a GAM (see Supplementary Material D for details). The bottom plot shows the number of recordings per 10-minute segment, showing the survey effort from 05:00–14:00. Peak activity occurs shortly after dawn, dropping rapidly but with some calling activity recorded throughout the morning. Plot inset shows the duration of independent call bouts detected by the classifier. Call bouts are intervals of regular calling, with no detected call 200s either side of the bout. Daily calling typically consists of a number of calling bouts.

Location	Survey days	% days calls detected	Mean calling time per day (min)	Mean start time of first bout	Mean end time of last bout
1	87	70	24.2	07:34	09:41
2	90	46	29.9	06:58	09:12
3	103	82	31.3	07:30	10:15
4	105	86	26.5	07:44	09:52
5	79	33	29.9	07:31	09:23
6	103	79	24.4	07:56	10:15
7	129	89	30.9	06:53	09:54
8	105	65	40.8	06:16	10:01

Table 2: Calling behaviour across 8 survey locations for the 161 day survey period March–August 2016. Recorders were situated within the known home ranges of the four Hainan gibbon social groups existing during the study period, at locations intermediate between known home ranges, and in an area where a solitary male gibbon was thought to occur. Locations of home ranges are indicated by numbers 1, 2, 3 and 4. 6 = solitary.

Where false negatives are particularly costly, this is easily incorporated by lowering the threshold required for manual verification. We expect that with more, and more diverse, training data, error rates would decline further.

Where environmental conditions were similar to those used to train the model, predictions were almost perfect and could be used to identify start and end times of call phrases and bouts, returning almost identical values to a human observer. It is impossible to know in advance whether environmental conditions are similar enough to warrant confidence in the associated predictions, but these results suggest that, as more training data covering a range of environmental conditions are added, model applications may go beyond gibbon detection, by automatically extracting inputs for more detailed behavioural analyses, for example of gibbon call syntax (Clarke, Reichard, & Zuberbühler, 2006).

Practically, developing an acoustic classifier such as ours requires a number of steps: deciding on an appropriate unit of analysis; manually labelling data; augmenting data and allocating it between training, validation, and test sets; choosing and fitting appropriate neural network models; and selecting a preferred model and using it to process the unlabelled portion of the data. Our study illustrates how model development and implementation are informed and guided by ecological objectives, here primarily detecting gibbon vocalizations over time scales of minutes or hours, and domain knowledge of Hainan gibbon call behaviour.

We based our classifier on phrases, rather than shorter notes or longer calling bouts, to balance ease of identification with data availability and computational requirements. Individual notes are easily confused with other sources (see Figure 2b). While calling bouts are highly distinctive, there are relatively few of them and, being longer in duration, they require more parameters to capture the same degree of detail. Phrases are far more numerous, less variable, and require fewer parameters.

Given this choice, segment duration was chosen to be longer than the longest phrase across all training data (8 seconds). The slightly longer segment length provides more presence segments – for example, an 8s phrase results in three 10s presence segments, but would only result in a single segment if the segment length was restricted to 8s. Preliminary runs based on shorter segments of 0.5–2 seconds and *partial* phrases did not yield good performance, with many false positives, probably because a small segment is not enough to distinguish gibbons from other species calling within the same frequency range.

Even using phrases, we have relatively few positive examples and these occur within a highly variable background environment, which is likely to be a common situation for ecological applications. The amount of data available to train neural networks is important, and CNNs tend to require relatively large amounts of data (at least thousands of each class) to generalize well. It may often be possible, as in our case, to collect or label additional data, but data augmentation is a valuable low-cost strategy for increasing sample sizes in conjunction with these other more

effort-intensive approaches (Bergler et al., 2019; Hestness et al., 2017; Kahl et al., 2017; Sun et al., 2017). In practice the process can be an iterative one guided by subjective judgement. We initially annotated only 40h across five recordings, but models based on these were poor, even with augmentation. Model performance (on the same test set) improved as we add more training data; we were also able to create more complex neural networks. Gains in accuracy decreased with additional annotations, and we stopped when these became marginal, but presumably further increases are possible as novel environments are included.

Training, validation and test datasets should be constructed by allocating longer contiguous sequences of audio to each of these, and then preprocessing each of these, rather than randomly allocating the segments themselves, which are highly autocorrelated and will thus overstate test accuracy. Wherever possible, we recommend using entirely independent recordings in the test dataset.

We found that 2-D CNNs based on spectrograms performed substantially better than 1-D CNNs that use amplitude time series following some initial preprocessing, mirroring Stowell, Wood, et al. (2019). Deep neural networks are often motivated by an argument that they learn salient features, rather than having to have these provided to them, but where intermediate features (here, spectral densities) can be provided, these speed up the learning process and provide measurable benefits. Beyond the 2-D/1-D distinction, we found that network architectures had relatively little impact on model accuracy, and we achieved good performance using relatively small, simple network architectures, again motivated by limitations on training data. We used few dense layers, each with only a small number of nodes, as these are particularly parameter hungry. Our basic approach was to start with simple architectures, evaluate them, and then add complexity in an iterative manner.

Traditional performance metrics such as precision and recall, while important, are not the only relevant measures of classifier success. Practically, classifiers such as ours can be used to point to audio segments that possibly contain gibbon calls, and that require manual verification. Where classification accuracy lags behind that of human experts, or where errors are costly – that is, in many ecological applications – attention shifts from replacing manual annotation to facilitating it. Probability cutoffs can be calibrated to balance the costs of false positives and negatives, and, even if the model is wrong by a few seconds, the amount of time spent in manual verification, compared to that required to processing the entire file manually, is minimal. Our classifier reduces an eight-hour recording to on average 22 minutes with false positive and negative rates under 2%. This time can be further reduced by playing back only those 10s segments that are predicted to contain phrases, although in our case the reduction in overall time was offset by the difficulty of manually verifying segments that are often not contiguous in time.

Analysis of our multi-month dataset demonstrated that gibbons could be detected regularly across all selected survey points, with call detection consistent with known patterns of gibbon behaviour and ecology. Calls were detected at expected times (Chan et al., 2005), and our dataset



provides a more precise baseline on Hainan gibbon call timing and duration. Hainan gibbon calling bouts were also generally detected less frequently during the wet season, a period when other gibbon species are also known to sing less frequently (Cheyne, 2008; Clink, Ahmad, & Klinck, 2020). Interestingly, call bouts recorded within the area occupied by a solitary male gibbon were amongst the shortest recorded bouts, and started and finished later than bouts from known social groups. While we cannot exclude the possibility of detecting group calls at this location, this finding suggests important new information on the behavioural ecology of solitary Hainan gibbons that may assist future monitoring and conservation planning.

It is uncertain whether within-recorder and between-recorder variation in calling bout detections represents variation in calling frequency between groups, and/or variation in detection effectiveness by recorders, with the latter possibility likely associated with specific recorder placement, local terrain, specific gibbon movement patterns across landscapes, and group home range size (cf. Bryant et al. (2017)). Future work could investigate detection likelihood in relation to specific environmental parameters and local weather conditions (e.g., rainfall, wind, temperature), data on which were not available for our survey period but are known to affect calling behaviour in other gibbons (Coudrat, Nanthavong, Ngoprasert, Suwanwaree, & Savini, 2015; Yin et al., 2016).

Where calls can be detected across multiple recording locations, acoustic spatial capture-recapture methods provide a means of estimating animal abundance (Stevenson et al., 2015). While our locations are too far apart for this to be feasible, this represents an important next step in monitoring a critically endangered population. Classifiers capable of discriminating between groups or individuals can be valuable inputs to this process (Augustine, Royle, Linden, & Fuller, 2020), as well as providing insight into the behavioural ecology of groups or individuals. We also recommend that call detection ranges should be determined for the specific field conditions at BNNR (e.g., slope, vegetation density), to calibrate monitoring effectiveness of specific recorders, and determine effective recorder placement (grid area/density) to ensure saturation of monitoring coverage. However, passive acoustic monitoring can now be introduced as an important component of the Hainan gibbon conservation toolkit, both for future use at BNNR and also to potentially detect unknown remnant gibbon populations elsewhere across Hainan (S. T. Turvey et al., 2017). Our classifier permits rapid and potentially real-time monitoring of Hainan gibbons, and we hope that the approach we describe in developing this classifier can serve as a roadmap for practitioners to implement their own classifier for other passive acoustic monitoring projects, and contribute to the effective conservation of calling species.

## Acknowledgements

We thank the Management Office of Bawangling National Nature Reserve for logistical assistance in the field. Fieldwork was funded by an Arcus Foundation grant to STT and a Wildlife Acoustics grant to JVB. ID is supported in part by funding from the National Research Foundation of South

393 Africa (Grant ID 90782, 105782). ED is supported by a postdoctoral fellowship from the African  
394 Institute for Mathematical Sciences South Africa, Stellenbosch University and the Next Einstein  
395 Initiative. This work was carried out with the aid of a grant from the International Development  
396 Research Centre, Ottawa, Canada ([www.idrc.ca](http://www.idrc.ca)), and with financial support from the Govern-  
397 ment of Canada, provided through Global Affairs Canada (GAC; [www.international.gc.ca](http://www.international.gc.ca)). We  
398 also thank the following rangers who contributed to data collection: Guang Wei, Zhong Zhao, Qing  
399 Lin, Jinbing Zhang, Zhicheng Zhang, Qianjin Li, Xiaoliang Fu, Zhengchong Zhou, Lubiao Huang,  
400 Zhengkun Ye, Zhenghai Zou, Jinqiang Wang, Wentao Han and Zengnan Xie.

## 401 Conflict of interest

402 The authors declare no competing interests.

## 403 Authors' contributions

404 ST, JB and HM conceived the passive monitoring project and developed the study designs and  
405 protocols. ED, ID and JH conceived the development of an automated classifier and designed the  
406 methodology. WL, ZL, QC, ZZ, HM and JB were responsible for fieldwork and data collection.  
407 ED, AH and JB annotated the data. ED constructed the classifier and performed the analysis.  
408 ED, ID and ST wrote the paper. All authors contributed critically to the drafts and gave final  
409 approval for publication.

## 410 Data accessibility

411 All code for training and testing the neural networks and conducting additional analyses is avail-  
412 able at <https://github.com/emmanueldufourq/GibbonClassifier> (Supplementary Material E).  
413 A subset of acoustic recordings, including training and testing labels, has been stored on Zenodo:  
414 <https://doi.org/10.5281/zenodo.3991714>.

## 415 References

- 416 Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., ... Zheng, X. (2015).  
417 *TensorFlow: Large-scale machine learning on heterogeneous systems*. Retrieved from  
418 <https://www.tensorflow.org/>
- 419 Augustine, B. C., Royle, J. A., Linden, D. W., & Fuller, A. K. (2020). Spatial proximity moderates  
420 genotype uncertainty in genetic tagging studies. *Proceedings of the National Academy of*  
421 *Sciences*, 117(30), 17903–17912.
- 422 Bergler, C., Schröter, H., Cheng, R. X., Barth, V., Weber, M., Nöth, E., ... Maier, A. (2019).  
423 Orca-spot: An automatic killer whale sound detection toolkit using deep learning. *Scientific*  
424 *Reports*, 9(1), 1–17.

425 Bermant, P. C., Bronstein, M. M., Wood, R. J., Gero, S., & Gruber, D. F. (2019). Deep machine  
426 learning techniques for the detection and classification of sperm whale bioacoustics. *Scientific*  
427 *Reports*, 9(1), 1–10.

428 Brockelman, W. Y., & Srikosamatara, S. (1993). Estimation of density of gibbon groups by use of  
429 loud songs. *American Journal of Primatology*, 29(2), 93–108. doi: 10.1002/ajp.1350290203

430 Bryant, J. V., Brulé, A., Wong, M. H., Hong, X., Zhou, Z., Han, W., ... Turvey, S. T. (2016).  
431 Detection of a new hainan gibbon (*nomascus hainanus*) group using acoustic call playback.  
432 *International Journal of Primatology*, 37(4-5), 534–547.

433 Bryant, J. V., Zeng, X., Hong, X., Chatterjee, H. J., & Turvey, S. T. (2017). Spatiotemporal  
434 requirements of the hainan gibbon: Does home range constrain recovery of the world’s rarest  
435 ape? *American Journal of Primatology*, 79(3), e22617.

436 Cannam, C., Landone, C., & Sandler, M. (2010, October). Sonic visualiser: An open source  
437 application for viewing, analysing, and annotating music audio files. In *Proceedings of the*  
438 *ACM Multimedia 2010 International Conference* (pp. 1467–1468). Firenze, Italy.

439 Chan, B. P. L., Fellowes, J., Geissmann, T., & Zhang, J. (2005). *Hainan gibbon status survey and*  
440 *conservation action plan. technical report 3*.

441 Cheyne, S. M. (2008). Effects of meteorology, astronomical variables, location and human distur-  
442 bance on the singing apes: *Hylobates albibarbis*. *American Journal of Primatology*, 70(4),  
443 386–392.

444 Chollet, F., et al. (2015). *Keras*. <https://keras.io>.

445 Christin, S., Hervet, E., & Lecomte, N. (2019). Applications for deep learning in ecology. *Methods*  
446 *in Ecology and Evolution*, 10(10), 1632–1644.

447 Clarke, E., Reichard, U. H., & Zuberbühler, K. (2006). The syntax and meaning of wild gibbon  
448 songs. *PloS one*, 1(1).

449 Clink, D. J., Ahmad, A. H., & Klinck, H. (2020). Gibbons aren’t singing in the rain: presence  
450 and amount of rainfall influences ape calling behavior in sabah, malaysia. *Scientific Reports*,  
451 10(1), 1–13.

452 Clink, D. J., Crofoot, M. C., & Marshall, A. J. (2019). Application of a semi-automated vocal finger-  
453 printing approach to monitor bornean gibbon females in an experimentally fragmented land-  
454 scape in sabah, malaysia. *Bioacoustics*, 28(3), 193–209. doi: 10.1080/09524622.2018.1426042

455 Coudrat, C., Nanthavong, C., Ngoprasert, D., Suwanwaree, P., & Savini, T. (2015). Singing  
456 patterns of white-cheeked gibbons (*nomascus* sp.) in the annamite mountains of laos. *Inter-*  
457 *national Journal of Primatology*, 36(4), 691–706.

458 Crunchant, A.-S., Borchers, D., Kuhl, H., & Piel, A. (2020). Listening and watching: Do cam-  
459 era traps or acoustic sensors more efficiently detect wild chimpanzees in an open habitat?  
460 *Methods in Ecology and Evolution*, 11(4), 542–552. doi: 10.1111/2041-210X.13362

461 Deng, H., Zhou, J., & Yang, Y. (2014). Sound spectrum characteristics of songs of hainan gibbon

(*nomascus hainanus*). *International Journal of Primatology*, 35(2), 547–556.

Fairbrass, A. J., Firman, M., Williams, C., Brostow, G. J., Titheridge, H., & Jones, K. E. (2019). CityNet-deep learning tools for urban ecoacoustic assessment. *Methods in Ecology and Evolution*, 10(2), 186–197.

Grill, T., & Schlüter, J. (2017). Two convolutional neural networks for bird detection in audio signals. , 1764–1768.

Hestness, J., Narang, S., Ardalani, N., Diamos, G., Jun, H., Kianinejad, H., ... Zhou, Y. (2017). Deep learning scaling is predictable, empirically. *Preprint* <https://arxiv.org/abs/1712.00409>.

Huang, X., Acero, A., & Hon, H.-W. (2001). *Spoken language processing: A guide to theory, algorithm, and system development*. Prentice Hall PTR.

Jiang, J.-j., Bu, L.-r., Duan, F.-j., Wang, X.-q., Liu, W., Sun, Z.-b., & Li, C.-y. (2019). Whistle detection and classification for whales based on convolutional neural networks. *Applied Acoustics*, 150, 169–178.

Kahl, S., Wilhelm-Stein, T., Hussein, H., Klinck, H., Kowerko, D., Ritter, M., & Eibl, M. (2017). Large-scale bird sound classification using convolutional neural networks. In *Conference and Labs of the Evaluation Forum*.

Kidney, D., Rawson, B. M., Borchers, D. L., Stevenson, B. C., Marques, T. A., & Thomas, L. (2016, 05). An efficient acoustic density estimation method with human detectors applied to gibbons in cambodia. *PLOS ONE*, 11(5), 1-16. doi: 10.1371/journal.pone.0155066

Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *Preprint* <https://arxiv.org/abs/1412.6980>.

Kiskin, I., Zilli, D., Li, Y., Sinka, M., Willis, K., & Roberts, S. (2020). Bioacoustic detection with wavelet-conditioned convolutional neural networks. *Neural Computing and Applications*, 32(4), 915–927.

Kulyukin, V., Mukherjee, S., & Amlathe, P. (2018). Toward audio beehive monitoring: Deep learning vs. standard machine learning in classifying beehive audio samples. *Applied Sciences*, 8(9), 1573.

LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436–444.

Liu, H., Ma, H., Cheyne, S. M., & Turvey, S. T. (2020). Recovery hopes for the world's rarest primate. *Science*, 368(6495), 1074–1074.

McFee, B., Lostanlen, V., McVicar, M., Metsai, A., Balke, S., Thome, C., ... Weiss, A. (2020). *Librosa*. doi: <https://doi.org/10.5281/zenodo.3606573>

Nolasco, I., Terenzi, A., Cecchi, S., Orcioni, S., Bear, H. L., & Benetos, E. (2019). Audio-based identification of beehive states. In *2019 International Conference on Acoustics, Speech and Signal Processing* (pp. 8256–8260).

Pamula, H., Pocha, A., & Klaczynski, M. (2019). Towards the acoustic monitoring of birds migrating at night. *Biodiversity Information Science and Standards*, 3, e36589.

499 Patterson, J., & Gibson, A. (2017). *Deep learning: a practitioner's approach*. O'Reilly Media,  
500 Inc.

501 Qazi, K. A., Tabassam Nawaz, Z. M., Rashid, M., & Habib, H. A. (2018). A hybrid technique for  
502 speech segregation and classification using a sophisticated deep neural network. *PloS One*,  
503 13(3).

504 Sethi, S. S., Jones, N. S., Fulcher, B. D., Picinali, L., Clink, D. J., Klinck, H., ... Ewers, R. M.  
505 (2020). Characterizing soundscapes across diverse ecosystems using a universal acoustic  
506 feature set. *Proceedings of the National Academy of Sciences*, 117(29), 17049–17055.

507 Shiu, Y., Palmer, K., Roch, M. A., Fleishman, E., Liu, X., Nosal, E.-M., ... Klinck, H. (2020).  
508 Deep neural networks for automated detection of marine mammal species. *Scientific Reports*,  
509 10(1), 1–12.

510 Sprengel, E., Jaggi, M., Kilcher, Y., & Hofmann, T. (2016). Audio based bird species identification  
511 using deep learning techniques. *2016 Conference and Labs of the Evaluation Forum*.

512 Stevenson, B. C., Borchers, D. L., Altwegg, R., Swift, R. J., Gillespie, D. M., & Measey, G. J.  
513 (2015). A general framework for animal density estimation from acoustic detections across a  
514 fixed microphone array. *Methods in Ecology and Evolution*, 6(1), 38–48.

515 Stowell, D., Petrusková, T., Šálek, M., & Linhart, P. (2019). Automatic acoustic identification of  
516 individuals in multiple species: improving identification across recording conditions. *Journal*  
517 *of the Royal Society Interface*, 16(153), 20180940.

518 Stowell, D., Wood, M. D., Pamuła, H., Stylianou, Y., & Glotin, H. (2019). Automatic acoustic  
519 detection of birds through deep learning: the first bird audio detection challenge. *Methods*  
520 *in Ecology and Evolution*, 10(3), 368–380.

521 Sun, C., Shrivastava, A., Singh, S., & Gupta, A. (2017). Revisiting Unreasonable Effectiveness of  
522 Data in Deep Learning Era. In *2017 International Conference on Computer Vision* (p. 843–  
523 852). doi: 10.1109/ICCV.2017.97

524 Turvey, S., Traylor-Holzer, K., Wong, M., Bryant, J., Zeng, X., Hong, X., & Long, Y. (2015).  
525 International conservation planning workshop for the hainan gibbon: final report. *Zoological*  
526 *Society of London, London, UK IUCN SSC Conservation Breeding Specialist Group, Apple*  
527 *Valley, MN, USA*.

528 Turvey, S. T., Bryant, J. V., Duncan, C., Wong, M. H., Guan, Z., Fei, H., ... Fan, P. (2017). How  
529 many remnant gibbon populations are left on hainan? Testing the use of local ecological  
530 knowledge to detect cryptic threatened primates. *American Journal of Primatology*, 79(2),  
531 e22593.

532 Vu, T. T., & Tran, L. M. (2019, Oct-01-2019). An application of autonomous recorders for gibbon  
533 monitoring. *International Journal of Primatology*. doi: 10.1007/s10764-018-0073-3

534 Wood, S. N. (2017). *Generalized additive models: an introduction with r*. CRC press.

535 Yin, L.-Y., Fei, H.-L., Chen, G.-S., Li, J.-H., Cui, L.-W., & Fan, P.-F. (2016). Effects of group

536 density, hunting, and temperature on the singing patterns of eastern hoolock gibbons (hoolock  
537 leuconedys) in gaoligongshan, southwest china. *American Journal of Primatology*, 78(8),  
538 861–871.

539 Zhang, H., Wang, C., Turvey, S. T., Sun, Z., Tan, Z., Yang, Q., ... Yang, D. (2020).  
540 Thermal infrared imaging from drones can detect individuals and nocturnal behavior  
541 of the world's rarest primate. *Global Ecology and Conservation*, 23, e01101. doi:  
542 <https://doi.org/10.1016/j.gecco.2020.e01101>

543 Zhou, X., Guan, Z., Zhong, E., Dong, Y., Li, H., & Hu, K. (2019). Automated monitoring of  
544 western black crested gibbon population based on voice characteristics. *5th International  
545 Conference on Computer and Communications*, 1383-1387.



## Supplementary Material

### A Details of observed call bouts in training data

In the preliminary stage of model building we used a subset of 72 hours of recordings (nine eight-hour recordings) to inform our decision to use a window of 10s. Across these recordings, an average of 2.3 calling bouts were observed per eight-hour period (min 1, max 4), with on average 54 phrases per bout (min 31, max 116). The average duration between phrases within a calling bout was 19.4s. Table A.1 presents the distribution of the numbers of syllables per phrase, as well as the mean duration of phrases consisting of different numbers of syllables. All phrases contained between one and six phrases, with the majority of phrases made up of one to four syllables.

Type	Average total duration	Number of instances
1 syllable	$2.6 \pm 0.9$	175
2 syllables	$4.3 \pm 0.8$	413
3 syllables	$5.1 \pm 0.9$	339
4 syllables	$5.8 \pm 0.8$	302
5 syllables	$6.4 \pm 0.8$	62
6 syllables	$6.0 \pm 0.6$	13
Duet	$6.1 \pm 0.9$	56

Table A.1: The average total duration for each type of hainan gibbon song. These are the syllables in the long calls that the hainan gibbon's perform. The number of times each type occurs is also presented. These values also include the breaks between consecutive calls.

### B Details of predicted call bouts in test data

Table B.1 shows observed and predicted start and end times of calling bouts in nine eight-hour recordings used to test our final (2-D CNN) model. Each bout is denoted by  $[t_s, t_e]$ , where  $t_s$  and  $t_e$  are start and end times (in seconds from the start of the recording) respectively. No calling bouts were missed, but two predicted bouts were false positives (denoted in bold) - these are 52 and 272 seconds of false positives respectively.

### C Seasonal differences in gibbon detections

Table C.1 reports the same summary statistics as Table 2 in the main text, but separately for wet and dry seasons. Gibbons called substantially less frequently in the wet season at four sites (2, 5, 8), more frequently at two sites (4, 7), and less on average across all sites (65% (261/403) vs. 77% (305/398)). Calling occurred over a substantially greater part of the day in the wet season (07:05–10:25) than in the dry season (07:25–09:30), although mean calling time per day did not differ substantially (wet season = 30m, dry season = 29m).

File	Type	Bouts (seconds)
1	Correct	[3682 3899], [3911 4174]
	Predicted	[3357, 4180]
2	Correct	[3349 3831], [8854 9456], [14796 15502]
	Predicted	[3342, 3836], [8850, 9459], [14791, 15506]
3	Correct	[3676 3795], [14759 14955], [19557 20257], [20533 20856]
	Predicted	[3623, 3802], [14752, 14962], [19365, 20262], [20526, 20860]
4	Correct	[3950 4201], [5390 5941]
	Predicted	[3945, 4208], [5351, 5948]
5	Correct	[3398 4148], [8507 9018], [10642 11035], [14918 15542]
	Predicted	[3366, 4154], [8477, 9024], [10509, 11039], [14911, 15548]
6	Correct	[3423 3783], [6370 7086]
	Predicted	[1216, 1268], [3417, 3789], [6367, 7091]
7	Correct	[5607 6626]
	Predicted	[1704, 1976], [5627, 6629]
8	Correct	[3133 3802], [11643 12317]
	Predicted	[3312, 4028], [11488, 12322]
9	Correct	[10210 10235], [24377 25125]
	Predicted	[10184, 10239], [24373, 25129]

Table B.1: Observed and predicted start and end times (sec) of calling bouts

Location	Survey days	% days calls detected	Mean calling time per day (min)	Mean start time of first bout	Mean end time of last bout
<b>Dry season (Mar-Apr), 61 days</b>					
1	52	75	28.7	07:19	09:26
2	37	73	26.4	07:00	08:48
3	58	81	34.8	08:01	10:09
4	52	81	28.6	07:46	09:25
5	51	51	29.9	07:31	09:23
6	55	80	28.3	08:04	10:10
7	46	78	21.6	07:07	08:46
8	47	94	31.2	06:33	09:26
<b>Wet season (May-Aug), 100 days</b>					
1	35	63	16.2	08:00	10:09
2	53	26	36.8	06:54	09:58
3	45	82	26.7	06:49	10:24
4	53	91	24.5	07:42	10:16
5	28	0	—	—	—
6	48	77	19.9	07:48	10:20
7	83	95	35.2	06:47	10:25
8	58	41	58.5	05:45	11:05

Table C.1: Detection of gibbon calling bouts by different recorders in wet and dry seasons.

## 568 D Generalized additive model details

569 We fitted a generalized additive model (GAM) with the *mgcv* package in R (Wood, 2017) to  
 570 model the relationship between the number of detected gibbon call bouts and time-of-day. A  
 571 binomial distribution for the error terms and an log link function was used, with a smooth term  
 572 using cubic regression splines with 10 knots ( $k = 10$ ) capturing non-linearities in the relationship  
 573 between predictor and response variable. The exact number of knots is not critical but was chosen  
 574 conservatively with the intention of producing biologically meaningful results. We checked that  
 575 we did not over-specify the number of knots using the effective degrees of freedom as a guide.  
 576 The model explained 88% of the variability in detected counts (deviance explained) and residual  
 577 analysis plots indicated symmetrically distributed residuals. There was no discernible evidence  
 578 of heteroskedasticity or unmodelled relationships between residuals and either observed or fitted  
 579 values of the dependent variable.

## 580 E Software

581 Two interactive notebooks, *Train.ipynb* and *Predict.ipynb*, illustrate the two main processes in  
 582 developing an automated classifier: pre-processing audio file and training a convolutional neural  
 583 network (*Train.ipynb*) and using an already-constructed model to identify calls in a new and  
 584 unlabelled recordings (*Predict.ipynb*).

585 A detailed manual is provided in the same repository as the code, so here we only briefly  
 586 illustrate the workflow (Figure E.1). Users first need to download the code repository and install  
 587 all requirements in *requirements.txt* using `pip install -r requirements.txt`.

588 For training, input data takes the form of (a) one or more .wav files containing already annotated  
 589 recordings, (b) a text file containing the annotated call times in the training files, and (c) a text file  
 590 containing the filenames of these .wav files. Upon downloading the repository, an example of (a)  
 591 and (c) is downloaded to the *Raw\_Data/Train* and *Call\_Labels* folders, while an example of (b)  
 592 appears as *Training\_Files.txt* in the root directory. Folders and filenames can be changed as note-  
 593 book options, as well as optional parameters controlling various aspects of model building (down-  
 594 sampling rate, augmentation, etc). Two core functions `execute_preprocessing_all_files` and  
 595 `train_model` perform preprocessing (creating image files containing spectrograms) and build the  
 596 CNNs.

597 For predicting on test or unlabelled files, users specify the location of the test .wav file, as well as  
 598 the location of the model parameters obtained during training. Model weights for our best 2-D CNN  
 599 are downloaded with the repository and saved as *Experiments/pretrained\_weights\_from\_paper.hdf5*,  
 600 so this notebook will run directly on new data without needing to retrain the model. The function  
 601 `execute_processing` runs the test file through the trained neural network and outputs predicted

602 call times as a spreadsheet.

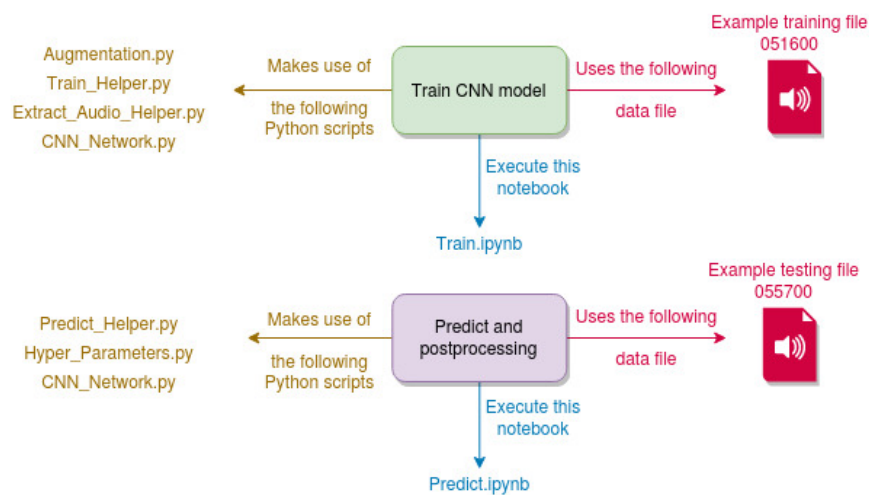


Figure E.1: Illustrating the pipeline and code dependencies for training and prediction.